

Import Important Libraries

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

Load Dataset

```
df = pd.read_csv("student_scores.csv")
```

Let's Analyze Dataset

Analyze attributes

```
df.head(3)
```

	Unnamed: 0	Gender	EthnicGroup	ParentEduc	LunchType
TestPrep \					
0	0	female	NaN	bachelor's degree	standard
1	1	female	group C	some college	standard
2	2	female	group B	master's degree	standard

	ParentMaritalStatus	PracticeSport	IsFirstChild	NrSiblings
TransportMeans \				
0	married	regularly	yes	3.0
1	married	sometimes	yes	0.0
2	single	sometimes	yes	4.0

	WklyStudyHours	MathScore	ReadingScore	WritingScore
0	< 5	71	71	74
1	5 - 10	69	90	88
2	< 5	87	93	91

```
df.tail()
```

	Unnamed: 0	Gender	EthnicGroup	ParentEduc
LunchType \				
30636	816	female	group D	high school
30637	890	male	group E	high school

```

standard
30638      911  female      NaN      high school
free/reduced
30639      934  female  group D  associate's degree
standard
30640      960   male  group B      some college
standard

      TestPrep ParentMaritalStatus PracticeSport IsFirstChild
NrSiblings \
30636      none              single      sometimes          no
2.0
30637      none              single      regularly          no
1.0
30638  completed              married      sometimes          no
1.0
30639  completed              married      regularly          no
3.0
30640      none              married      never              no
1.0

      TransportMeans WklyStudyHours  MathScore  ReadingScore
WritingScore
30636      school_bus      5 - 10      59      61
65
30637      private      5 - 10      58      53
51
30638      private      5 - 10      61      70
67
30639      school_bus      5 - 10      82      90
93
30640      school_bus      5 - 10      64      60
58

```

Check how many rows and columns are present in dataset

```

df.shape

(30641, 15)

```

Get info of dataset

```

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30641 entries, 0 to 30640
Data columns (total 15 columns):
 #   Column              Non-Null Count  Dtype
---  -

```

```

0  Unnamed: 0      30641 non-null int64
1  Gender          30641 non-null object
2  EthnicGroup     28801 non-null object
3  ParentEduc      28796 non-null object
4  LunchType       30641 non-null object
5  TestPrep        28811 non-null object
6  ParentMaritalStatus 29451 non-null object
7  PracticeSport   30010 non-null object
8  IsFirstChild    29737 non-null object
9  NrSiblings      29069 non-null float64
10 TransportMeans  27507 non-null object
11 WklyStudyHours  29686 non-null object
12 MathScore       30641 non-null int64
13 ReadingScore    30641 non-null int64
14 WritingScore    30641 non-null int64
dtypes: float64(1), int64(4), object(10)
memory usage: 3.5+ MB

```

Analyze columns one by one

```

df['Gender'].unique()
array(['female', 'male'], dtype=object)

df['EthnicGroup'].unique()
array([nan, 'group C', 'group B', 'group A', 'group D', 'group E'],
      dtype=object)

```

We see that there are some missing values present in EthnicGroup column

```

df['ParentEduc'].unique()
array(["bachelor's degree", 'some college', "master's degree",
      "associate's degree", 'high school', 'some high school', nan],
      dtype=object)

df['TransportMeans'].unique()
array(['school_bus', nan, 'private'], dtype=object)

df['WklyStudyHours'].unique()
array(['< 5', '5 - 10', '> 10', nan], dtype=object)

```

Describe data

```
df.describe()
```

	Unnamed: 0	NrSiblings	MathScore	ReadingScore	WritingScore
count	30641.000000	29069.000000	30641.000000	30641.000000	30641.000000
mean	499.556607	2.145894	66.558402	69.377533	68.418622
std	288.747894	1.458242	15.361616	14.758952	15.443525
min	0.000000	0.000000	0.000000	10.000000	4.000000
25%	249.000000	1.000000	56.000000	59.000000	58.000000
50%	500.000000	2.000000	67.000000	70.000000	69.000000
75%	750.000000	3.000000	78.000000	80.000000	79.000000
max	999.000000	7.000000	100.000000	100.000000	100.000000


```
df.isna().sum() ## Check how many values are null
```

Unnamed: 0	0
Gender	0
EthnicGroup	1840
ParentEduc	1845
LunchType	0
TestPrep	1830
ParentMaritalStatus	1190
PracticeSport	631
IsFirstChild	904
NrSiblings	1572
TransportMeans	3134
WklyStudyHours	955
MathScore	0
ReadingScore	0
WritingScore	0
dtype: int64	

Heatmap of missing values


```
plt.figure(figsize=(12, 8))
sns.heatmap(df.isnull(), cbar=False, cmap='viridis')
```

```
plt.title('Missing Data Heatmap')
plt.show()
```



Let's fill and drop values

```
plt.figure(figsize=(10,3))
sns.countplot(x=df['ParentEduc'])
plt.title("Count Plot of Parent's Education")
plt.show()
```




```
df['ParentEduc'].value_counts()
```

```
ParentEduc
some college    6633
high school    5687
associate's degree  5550
some high school  5517
bachelor's degree  3386
master's degree   2023
Name: count, dtype: int64
```

Let's check degree on the basis of Parent Marital Status

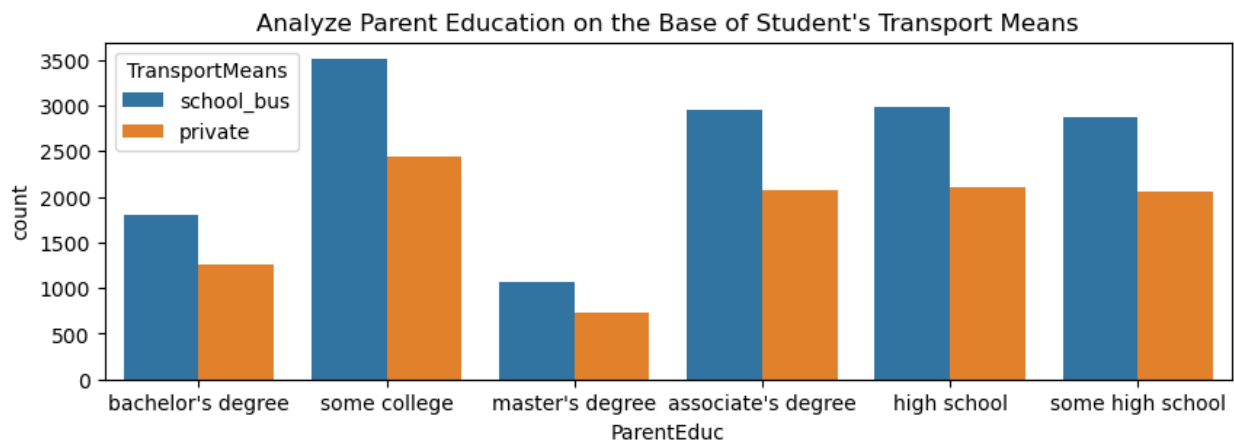
```
plt.figure(figsize=(10,3))
sns.countplot(x=df['ParentEduc'], hue=df['ParentMaritalStatus'])
plt.title("Analyze Parent Education on the Base of their Marital Status")
plt.show()
```



Analyze education on the base of transport means

```
plt.figure(figsize=(10,3))
sns.countplot(x=df['ParentEduc'], hue=df['TransportMeans'])
```

```
plt.title("Analyze Parent Education on the Base of Student's Transport Means")
plt.show()
```



As, I don't find any specific trend in Parent Educ with other, so the better option is dropna values from it

```
df.dropna(subset=['ParentEduc'], inplace=True)
```

```
df.isna().sum()
```

```
Unnamed: 0      0
Gender          0
EthnicGroup    1726
ParentEduc      0
LunchType      0
TestPrep      1722
ParentMaritalStatus  1116
PracticeSport   581
IsFirstChild   844
NrSiblings    1477
TransportMeans  2933
WklyStudyHours  893
MathScore      0
ReadingScore   0
WritingScore   0
dtype: int64
```

Let's first drop UnNamed Column from our dataset b/c it makes no sense

```
df.head(1)
```

```

    Unnamed: 0  Gender EthnicGroup      ParentEduc LunchType
TestPrep \
0          0  female          NaN  bachelor's degree  standard
none

    ParentMaritalStatus PracticeSport IsFirstChild  NrSiblings
TransportMeans \
0          married      regularly          yes          3.0
school_bus

    WklyStudyHours  MathScore  ReadingScore  WritingScore
0          < 5          71          71          74

df.drop('Unnamed: 0', axis=1, inplace=True)
# Axis = 1 specifies that i want to drop column, not row

df.head(1)

    Gender EthnicGroup      ParentEduc LunchType TestPrep \
0  female          NaN  bachelor's degree  standard  none

    ParentMaritalStatus PracticeSport IsFirstChild  NrSiblings
TransportMeans \
0          married      regularly          yes          3.0
school_bus

    WklyStudyHours  MathScore  ReadingScore  WritingScore
0          < 5          71          71          74

```

Analyze no. of siblings

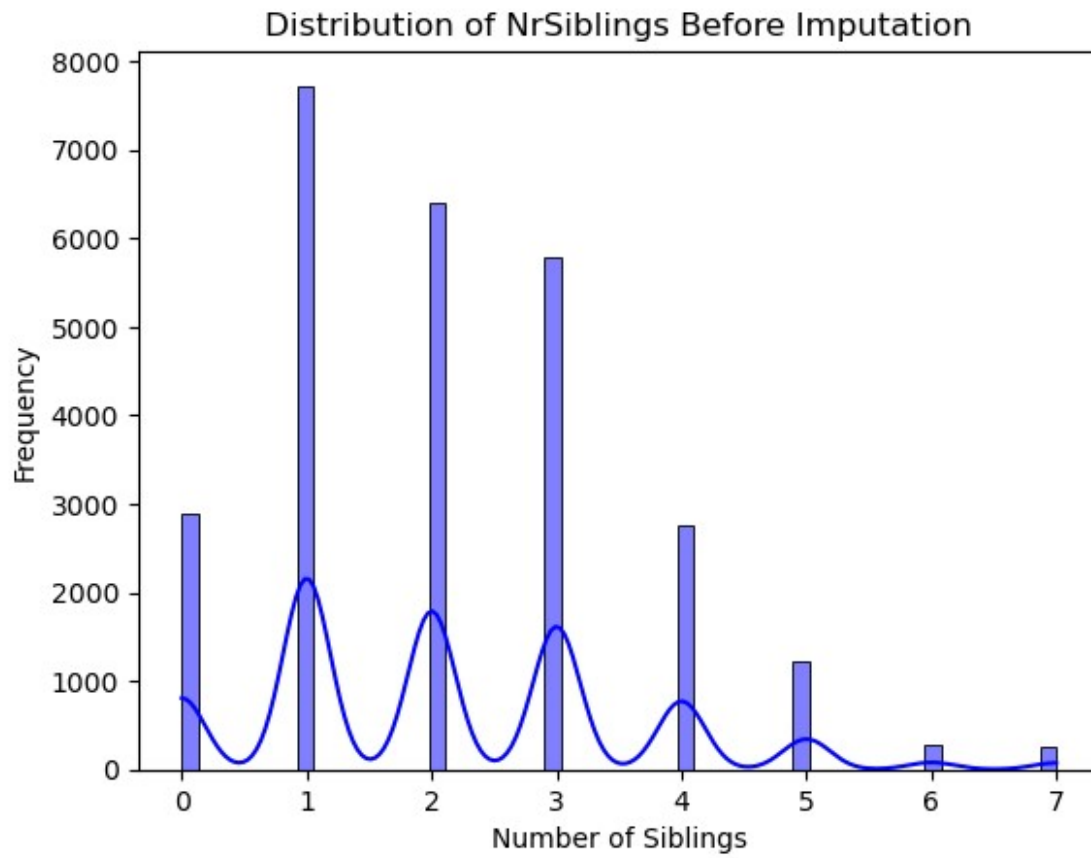
```

df['NrSiblings'].unique()

array([ 3.,  0.,  4.,  1., nan,  2.,  5.,  7.,  6.])

sns.histplot(df['NrSiblings'], kde=True, color='blue')
plt.title('Distribution of NrSiblings Before Imputation')
plt.xlabel('Number of Siblings')
plt.ylabel('Frequency')
plt.show()

```


```
## Fill Missing value with mean
df_filled = df.copy() # Create a copy to compare before and after
df_filled['NrSiblings'].fillna(df['NrSiblings'].mode())
```

```
0      3.0
1      0.0
2      4.0
3      1.0
4      0.0
```


```
...
30636  2.0
30637  1.0
30638  1.0
30639  3.0
30640  1.0
```

```
Name: NrSiblings, Length: 28796, dtype: float64
```

```
sns.histplot(df_filled['NrSiblings'], kde=True, color='green')
plt.title('Distribution of NrSiblings After Imputation')
plt.xlabel('Number of Siblings')
plt.ylabel('Frequency')
plt.show()
```



```
sns.histplot(df['NrSiblings'], kde=True, color='blue', label='Before Imputation')
sns.histplot(df_filled['NrSiblings'], kde=True, color='green', label='After Imputation', alpha=0.5)
plt.title('Comparison of NrSiblings Before and After Imputation')
plt.xlabel('Number of Siblings')
plt.ylabel('Frequency')
plt.legend()
plt.show()
```



As we see that it does not affect our data, if we fill no of sibling with mode, so we fill it in our original dataset

```
df.isna().sum()
```

```
Gender          0
EthnicGroup     1726
ParentEduc      0
LunchType       0
TestPrep        1722
ParentMaritalStatus  1116
PracticeSport    581
IsFirstChild     844
NrSiblings      1477
TransportMeans   2933
WklyStudyHours   893
MathScore        0
ReadingScore     0
WritingScore     0
dtype: int64
```

```
# Compute the mode of 'NrSiblings' (mode returns a Series, take the first value)
```

```

mode_value = df['NrSiblings'].mode()[0]

# Fill missing values with the mode
df['NrSiblings'] = df['NrSiblings'].fillna(mode_value)

df.isna().sum()

```

Gender	0
EthnicGroup	1726
ParentEduc	0
LunchType	0
TestPrep	1722
ParentMaritalStatus	1116
PracticeSport	581
IsFirstChild	844
NrSiblings	0
TransportMeans	2933
WklyStudyHours	893
MathScore	0
ReadingScore	0
WritingScore	0
dtype: int64	

```

df['NrSiblings'].unique()

array([3., 0., 4., 1., 2., 5., 7., 6.])

```


Compute the correlation matrix

```

corr = df[['MathScore', 'ReadingScore', 'WritingScore']].corr()

# Create a heatmap of the correlation matrix
plt.figure(figsize=(8, 6))
sns.heatmap(corr, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix of Scores')
plt.show()

```



Distribution of Score

```
plt.figure(figsize=(18, 6))


# Histogram for MathScore
plt.subplot(1, 3, 1)
sns.histplot(df['MathScore'], kde=True, color='blue')
plt.title('Distribution of Math Scores')
plt.xlabel('Math Score')
plt.ylabel('Frequency')

# Histogram for ReadingScore
plt.subplot(1, 3, 2)
sns.histplot(df['ReadingScore'], kde=True, color='green')
plt.title('Distribution of Reading Scores')
```

```
plt.xlabel('Reading Score')
plt.ylabel('Frequency')

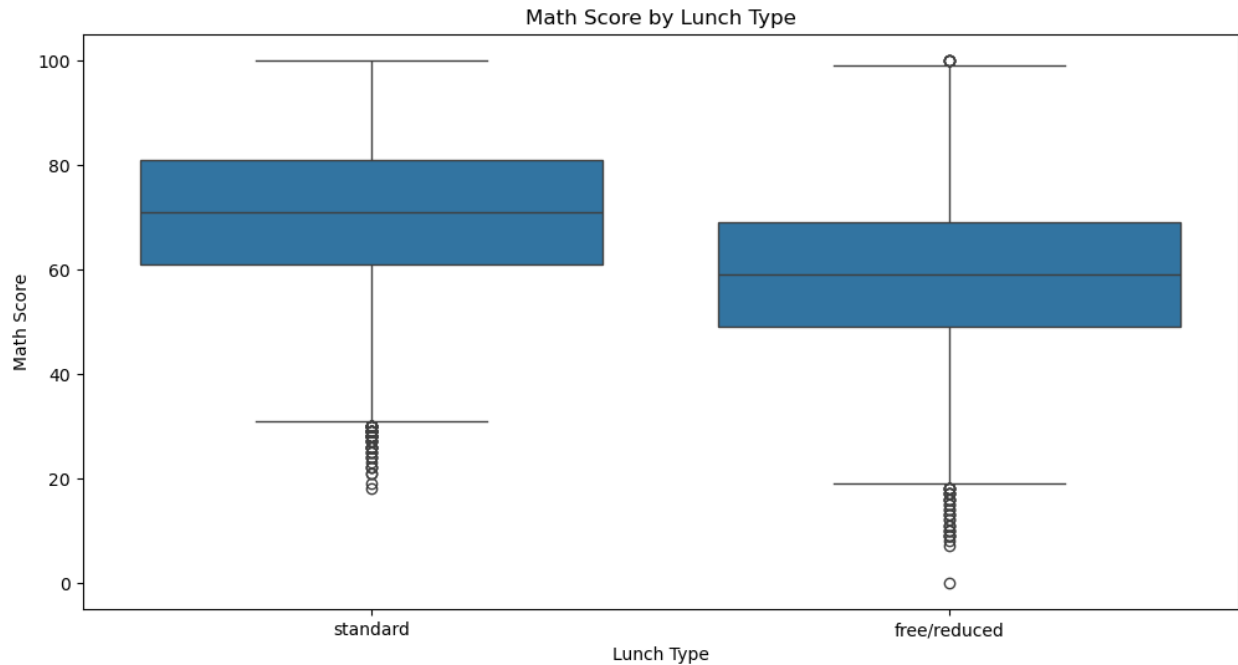
# Histogram for WritingScore
plt.subplot(1, 3, 3)
sns.histplot(df['WritingScore'], kde=True, color='red')
plt.title('Distribution of Writing Scores')
plt.xlabel('Writing Score')
plt.ylabel('Frequency')

plt.tight_layout()
plt.show()
```



Box plot for lunch type

```
# Box plot for MathScore by LunchType
plt.figure(figsize=(12, 6))
sns.boxplot(x='LunchType', y='MathScore', data=df)
plt.title('Math Score by Lunch Type')
plt.xlabel('Lunch Type')
plt.ylabel('Math Score')
plt.show()
```




Analyze Ethnic Group

```
df['EthnicGroup'].unique()

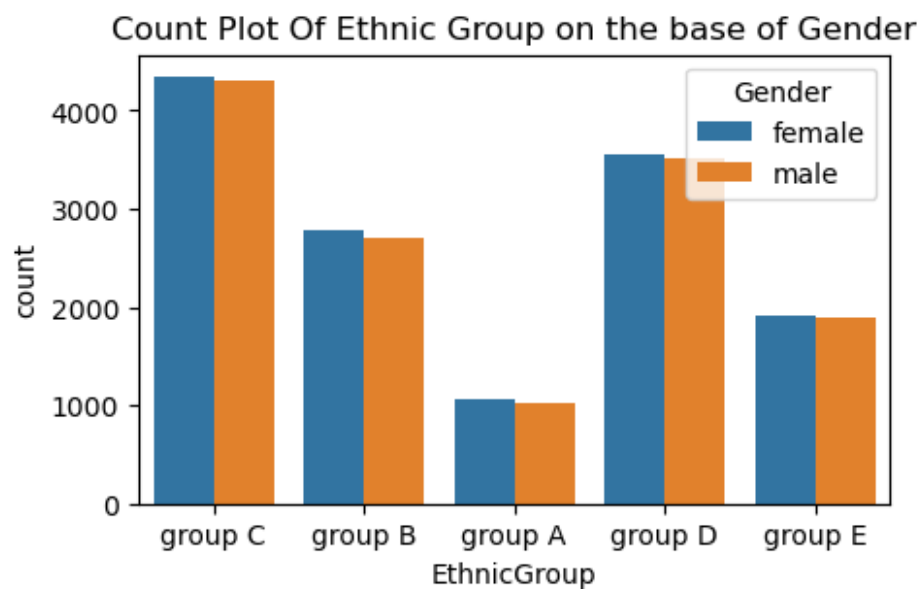
array([nan, 'group C', 'group B', 'group A', 'group D', 'group E'],
      dtype=object)

plt.figure(figsize=(6,4))
plt.title("Count Plot of Ethnic Group")
sns.countplot(df['EthnicGroup'])
plt.show()
```




Analyze Ethnic Group on the base of gender

```
plt.figure(figsize=(5,3))  
sns.countplot(x=df['EthnicGroup'], hue=df['Gender'])  
plt.title("Count Plot Of Ethnic Group on the base of Gender")  
plt.show()
```



Let's Analyze it with practice sport

```
plt.figure(figsize=(5,3))
sns.countplot(x=df['EthnicGroup'], hue=df['PracticeSport'])
plt.title("Count Plot of Ethnic Group on the base of Practice Sport")
plt.show()
```



It follow no trend, lets dropna

```
df.dropna(subset=['EthnicGroup'], inplace=True)
```

```
df.isna().sum()
```

```
Gender          0
EthnicGroup     0
ParentEduc      0
LunchType       0
TestPrep       1611
ParentMaritalStatus 1038
PracticeSport    535
IsFirstChild    796
NrSiblings       0
TransportMeans  2763
WklyStudyHours   829
MathScore        0
ReadingScore     0
WritingScore     0
dtype: int64
```


Lets drop all null values

```
df.dropna(inplace=True)
df.info()

<class 'pandas.core.frame.DataFrame'>
Index: 20266 entries, 2 to 30640
Data columns (total 14 columns):
 #   Column                Non-Null Count  Dtype
---  ---
 0   Gender                20266 non-null  object
 1   EthnicGroup           20266 non-null  object
 2   ParentEduc            20266 non-null  object
 3   LunchType             20266 non-null  object
 4   TestPrep              20266 non-null  object
 5   ParentMaritalStatus   20266 non-null  object
 6   PracticeSport         20266 non-null  object
 7   IsFirstChild          20266 non-null  object
 8   NrSiblings            20266 non-null  float64
 9   TransportMeans        20266 non-null  object
10   WklyStudyHours        20266 non-null  object
11   MathScore             20266 non-null  int64
12   ReadingScore          20266 non-null  int64
13   WritingScore          20266 non-null  int64
dtypes: float64(1), int64(3), object(10)
memory usage: 2.3+ MB
```

Analyze Gender distribution

```
plt.figure(figsize=(5,4))
ax = sns.countplot(x=df['Gender'])
ax.bar_label(ax.containers[0])
plt.title("Analyze Gender Distribution in School")
plt.show()
```



```
df['Gender'].value_counts()
```

```
Gender
female    10294
male      9972
Name: count, dtype: int64
```

Analysis

From the above chart, we analyze that no of female in a school is greater than no. on male students

Parent Education Impact on scores of Student


```
## We find the aggregate/estimate of parent education on student score
gb =
df.groupby('ParentEduc').agg({'MathScore':'mean', 'ReadingScore':'mean'
, 'WritingScore':'mean'})
gb
```

	MathScore	ReadingScore	WritingScore
ParentEduc			
associate's degree	68.483472	71.302296	70.496089

bachelor's degree	70.600165	73.383726	73.703428
high school	64.364892	67.354847	65.541688
master's degree	72.223796	75.825779	76.331445
some college	66.530300	69.230537	68.586802
some high school	62.558556	65.503246	63.645547

We use heatmap for this type of data analysis

```
plt.figure(figsize=(5,4))
sns.heatmap(gb, annot=True)
plt.title("Affect of Parents Education on Students Score")
plt.show()
```



Parents education impacts alot on their children scores

Parent Marital Status Impact on scores of Student

```
## We find the aggregate/estimate of parent education on student score
gb =
df.groupby('ParentMaritalStatus').agg({'MathScore': 'mean', 'ReadingScore': 'mean', 'WritingScore': 'mean'})
gb
```


	MathScore	ReadingScore	WritingScore
ParentMaritalStatus			

```

divorced      66.619941    69.742522    68.872434
married       66.675799    69.447144    68.455070
single        66.357481    69.428689    68.569665
widowed       67.966667    70.594872    69.543590

plt.figure(figsize=(5,4))
sns.heatmap(gb, annot=True)
plt.title("Effect of Parents Marital Status on Students Score")
plt.show()

```



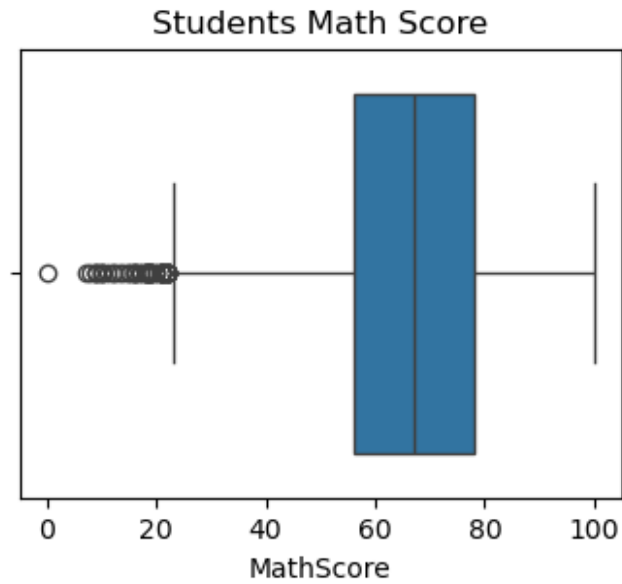
From the above chart we have concluded that there is no/negliable impact on students score due to their parents merital status

Lets detect outlier in scores

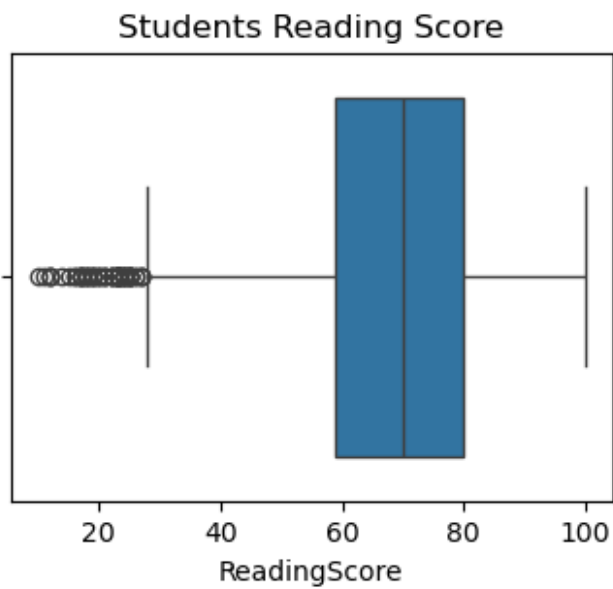
```

plt.figure(figsize=(4,3))
sns.boxplot(x=df['MathScore'])
plt.title("Students Math Score")
plt.show()


```



```
plt.figure(figsize=(4,3))
sns.boxplot(x=df['ReadingScore'])
plt.title("Students Reading Score")
plt.show()
```




```
plt.figure(figsize=(4,3))
sns.boxplot(x=df['WritingScore'])
plt.title("Students Writing Score")
plt.show()
```



Ethnics Group Distribution

```
values_count = df['EthnicGroup'].value_counts()

plt.figure(figsize=(5,3))
plt.title("Distribution Of Ethnic Group")
plt.pie(values_count, labels=values_count.keys(), autopct='%.2f%%')
plt.show()
```



```
values_count
```

```


EthnicGroup
group C    6484
group D    5252
group B    4116
group E    2862
group A    1552
Name: count, dtype: int64

values_count.keys()

Index(['group C', 'group D', 'group B', 'group E', 'group A'],
      dtype='object', name='EthnicGroup')

# Box plot for MathScore by EthnicGroup
plt.figure(figsize=(12, 6))
sns.boxplot(x='EthnicGroup', y='MathScore', data=df)
plt.title('Math Score by Ethnic Group')
plt.xlabel('Ethnic Group')
plt.ylabel('Math Score')
plt.show()

```




Effect of Gender on Score

Math Score


```
plt.figure(figsize=(12, 6))

# Box plot for MathScore by Gender
sns.boxplot(x='Gender', y='MathScore', data=df)
plt.title('Distribution of Math Scores by Gender')
plt.xlabel('Gender')
plt.ylabel('Math Score')
plt.show()
```



```
plt.figure(figsize=(12, 6))

# Violin plot for MathScore by Gender
sns.violinplot(x='Gender', y='MathScore', data=df)
plt.title('Distribution of Math Scores by Gender')
plt.xlabel('Gender')
plt.ylabel('Math Score')
plt.show()
```




```
# Summary statistics for MathScore by Gender
summary = df.groupby('Gender')['MathScore'].agg(['mean', 'median',
'std', 'count'])
print(summary)
```

	mean	median	std	count
Gender				
female	64.234991	64.0	15.142134	10294
male	69.071701	69.0	15.230742	9972

Reading Score

```
plt.figure(figsize=(12, 6))

# Box plot for Reading by Gender
sns.boxplot(x='Gender', y='ReadingScore', data=df)
plt.title('Distribution of Reading Scores by Gender')
plt.xlabel('Gender')
plt.ylabel('Math Score')
plt.show()
```



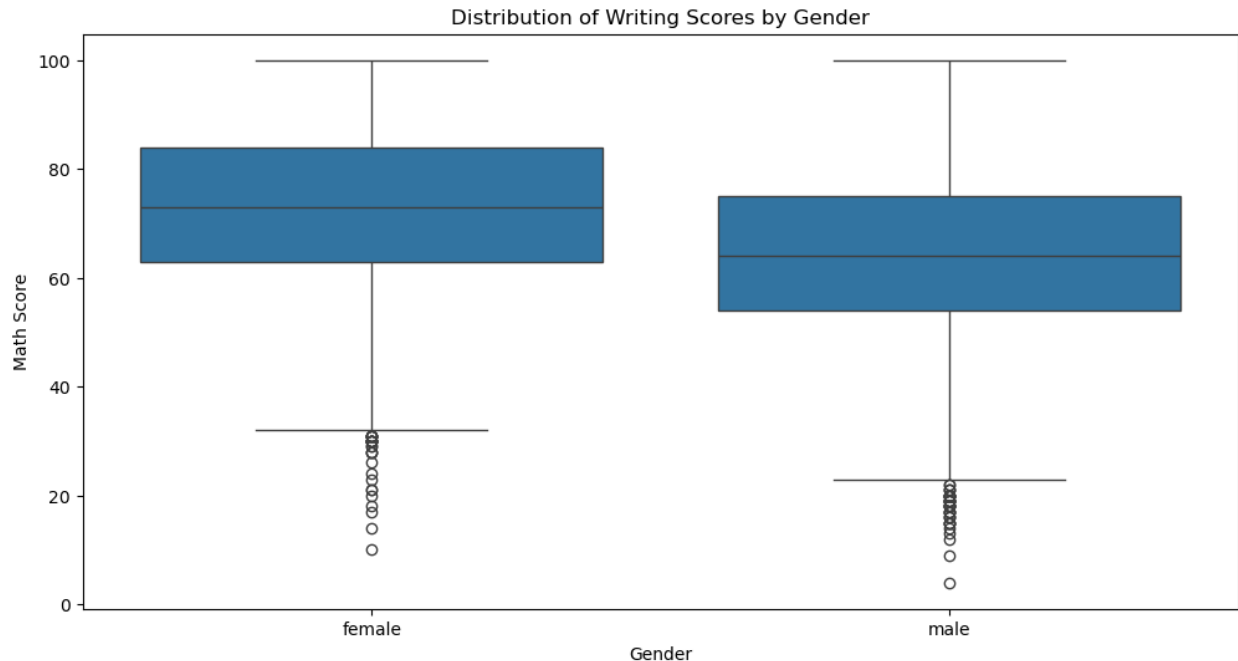
```
# Summary statistics for Reading by Gender
summary = df.groupby('Gender')['ReadingScore'].agg(['mean', 'median',
'std', 'count'])
print(summary)
```

	mean	median	std	count
Gender				
female	73.052943	73.0	14.153673	10294
male	65.861813	66.0	14.550282	9972

Writing Score

```
plt.figure(figsize=(12, 6))

# Box plot for Writing by Gender
sns.boxplot(x='Gender', y='WritingScore', data=df)
plt.title('Distribution of Writing Scores by Gender')
plt.xlabel('Gender')
plt.ylabel('Math Score')
plt.show()
```



```
# Summary statistics for Writing by Gender
summary = df.groupby('Gender')['WritingScore'].agg(['mean', 'median',
'std', 'count'])
print(summary)
```

	mean	median	std	count
Gender				
female	73.046338	73.0	14.572488	10294
male	63.956679	64.0	15.041557	9972

Effect of Weekly Study Hours on Student Score

```
# Set the size of the plots
plt.figure(figsize=(19, 6))


# Plot MathScore by TransportMeans
plt.subplot(1, 3, 1)
sns.boxplot(x='TransportMeans', y='MathScore', data=df)
plt.title('Math Score by Transport Means')
plt.xlabel('Transport Means')
plt.ylabel('Math Score')
plt.xticks(rotation=45)

# Plot ReadingScore by TransportMeans
```

```
plt.subplot(1, 3, 2)
sns.boxplot(x='TransportMeans', y='ReadingScore', data=df)
plt.title('Reading Score by Transport Means')
plt.xlabel('Transport Means')
plt.ylabel('Reading Score')
plt.xticks(rotation=45)

# Plot WritingScore by TransportMeans
plt.subplot(1, 3, 3)
sns.boxplot(x='TransportMeans', y='WritingScore', data=df)
plt.title('Writing Score by Transport Means')
plt.xlabel('Transport Means')
plt.ylabel('Writing Score')
plt.xticks(rotation=45)

plt.tight_layout()
plt.show()
```




```
plt.figure(figsize=(18, 6))

# Plot MathScore by TransportMeans
plt.subplot(1, 3, 1)
sns.violinplot(x='TransportMeans', y='MathScore', data=df)
plt.title('Math Score by Transport Means')
plt.xlabel('Transport Means')
plt.ylabel('Math Score')
plt.xticks(rotation=45)

# Plot ReadingScore by TransportMeans
plt.subplot(1, 3, 2)
sns.violinplot(x='TransportMeans', y='ReadingScore', data=df)
plt.title('Reading Score by Transport Means')
plt.xlabel('Transport Means')
plt.ylabel('Reading Score')
plt.xticks(rotation=45)
```

```
# Plot WritingScore by TransportMeans
plt.subplot(1, 3, 3)
sns.violinplot(x='TransportMeans', y='WritingScore', data=df)
plt.title('Writing Score by Transport Means')
plt.xlabel('Transport Means')
plt.ylabel('Writing Score')
plt.xticks(rotation=45)

plt.tight_layout()
plt.show()
```



```
# Summary statistics for MathScore by TransportMeans
summary_math = df.groupby('TransportMeans')['MathScore'].agg(['mean',
'median', 'std', 'count'])
print("Math Score Summary by Transport Means:\n", summary_math)

# Summary statistics for ReadingScore by TransportMeans
summary_reading = df.groupby('TransportMeans')
['ReadingScore'].agg(['mean', 'median', 'std', 'count'])
print("Reading Score Summary by Transport Means:\n", summary_reading)

# Summary statistics for WritingScore by TransportMeans
summary_writing = df.groupby('TransportMeans')
['WritingScore'].agg(['mean', 'median', 'std', 'count'])
print("Writing Score Summary by Transport Means:\n", summary_writing)
```

```
Math Score Summary by Transport Means:
              mean  median      std  count
TransportMeans
private      66.570783    67.0  15.608412   8406
school_bus   66.646206    67.0  15.210997  11860
Reading Score Summary by Transport Means:
              mean  median      std  count
TransportMeans
```

private	69.564478	70.0	14.951528	8406
school_bus	69.479089	70.0	14.680749	11860

Writing Score Summary by Transport Means:

	mean	median	std	count
TransportMeans				
private	68.629550	69.0	15.671928	8406
school_bus	68.534148	69.0	15.354334	11860

Effect of Test Preperation

```
# Box plot for MathScore by TestPrep
plt.figure(figsize=(12, 6))
sns.boxplot(x='TestPrep', y='MathScore', data=df)
plt.title('Math Score by Test Preparation')
plt.xlabel('Test Prep')
plt.ylabel('Math Score')
plt.show()
```

